

# 10-315 Introduction to Machine Learning (SCS Majors)

## Lecture 7: Logistic Regression

Leila Wehbe  
Carnegie Mellon University  
Machine Learning Department

Reading: <http://www.cs.cmu.edu/~tom/mlbook/NBayesLogReg.pdf>  
(<http://www.cs.cmu.edu/~tom/mlbook/NBayesLogReg.pdf>). Generative and Discriminative Classifiers by Tom Mitchell.

### Lecture outcomes:

- Logistic Regression
- Gradient Descent Review
- Comparing LR and GNB

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from scipy.stats import norm
import seaborn as sns
sns.set_theme()

x1 = np.linspace(-10,10,1000)
x2 = np.linspace(-10,10,1000)
```

## Naïve Bayes is a *Generative* classifier

Generative classifiers:

- Assume a functional form for  $P(X, Y)$  (or  $P(X|Y)$  and  $P(Y)$ )
- we can view  $P(X|Y)$  as describing how to sample random instances  $X$  given  $Y$ .

Instead of learning  $P(X|Y)$ , can we learn  $P(Y|X)$  directly or learn the decision boundary directly?

## *Discriminative* classifiers

- Assume some functional form for  $P(Y|X)$  or for the decision boundary
- Estimate parameters of  $P(Y|X)$  directly from training data

# Logistic Regression is a discriminative classifier

Learns  $f : X \rightarrow Y$ , where

- $X$  is a vector of real-valued or discrete features,  $(X_1, \dots, X_d)$
- $Y$  is boolean (can also be extended for  $K$  discrete classes).

$P(Y|X)$  is modeled as:


$$P(Y = 1|X) = \frac{1}{1 + \exp(-(w_0 + \sum_i w_i X_i))} = \frac{\exp(w_0 + \sum_i w_i X_i)}{\exp(w_0 + \sum_i w_i X_i) + 1}$$

It uses the logistic (or sigmoid) function:

$$\frac{1}{1 + \exp -z}$$

```
In [2]: z = np.linspace(-10,10,1000)
plt.plot(z,1/(1+np.exp(-z)))
plt.xlabel(r'$w_0+\sum_i w_i X_i$',fontsize=16)
plt.ylabel(r'$P(Y=1|X)$',fontsize=16)
```

Out[2]: Text(0, 0.5, '\$P(Y=1|X)\$')



# What is the form of the decision boundary?

$$\frac{P(Y = 1|X)}{P(Y = 0|X)} = \frac{\frac{\exp(w_0 + \sum_i w_i X_i)}{\exp(w_0 + \sum_i w_i X_i) + 1}}{\frac{1}{\exp(w_0 + \sum_i w_i X_i) + 1}} = \exp(w_0 + \sum_i w_i X_i)$$

Asking  $P(Y = 1|X) > P(Y = 0|X)$  is the same as asking if  $\ln \frac{P(Y=1|X)}{P(Y=0|X)} > 0$ .

i.e. is

$$w_0 + \sum_i w_i X_i > 0?$$


This is a linear decision boundary!

```
In [3]: from scipy.stats import multivariate_normal
# similar to previous example
mu_1_1 = -4; sigma_1_1 = 2; mu_2_1 = 4; sigma_2_1 = 2
mu_1_0 = 4; sigma_1_0 = 2; mu_2_0 = -4; sigma_2_0 = 2
cov_positive = np.array([[sigma_1_1**2, 3], [3, sigma_2_1**2]] )
cov_negative = np.array([[sigma_1_0**2, 3], [3, sigma_2_0**2]] )
# Sample data from these distributions
X_positive = multivariate_normal.rvs(mean=[mu_1_1, mu_2_1], cov=cov_positive, size
= (20))
X_negative = multivariate_normal.rvs(mean=[mu_1_0, mu_2_0], cov=cov_negative, size
= (20))
```

```
In [4]: plt.figure(figsize=(6,6))
plt.scatter(X_positive[:, 0], X_positive[:, 1],facecolors='r', edgecolors='w')
plt.scatter(X_negative[:, 0], X_negative[:, 1],facecolors='b', edgecolors='w')

# hand picked line
plt.plot(x1, x1*0.8+0.5)
from labellines import labelLine
labelLine(plt.gca().get_lines()[-1],0.6,label=r'$w_0+\sum_i w_i X_i = 0$',fontsize
=16)

plt.axis([-10,10,-10,10],'equal')
plt.xlabel(r'$x_1$',fontsize=20); plt.ylabel(r'$x_2$',fontsize=20)
plt.title('Data space',fontsize=20);
```



## Logistic Regression is a Linear Classifier

$$P(Y = 1|X) = \frac{\exp(w_0 + \sum_i w_i X_i)}{\exp(w_0 + \sum_i w_i X_i) + 1}$$
$$P(Y = 1|X) = \frac{1}{\exp(w_0 + \sum_i w_i X_i) + 1}$$

The weights  $w_i$  are optimized such that when  $w_0 + \sum_i w_i X_i > 0$  the example is more likely to be positive and when  $w_0 + \sum_i w_i X_i < 0$  it's more likely to be negative.



$$w_0 + \sum_i w_i X_i = 0, P(Y = 1|X) = \frac{1}{2}$$

$$w_0 + \sum_i w_i X_i \rightarrow \infty, P(Y = 1|X) \rightarrow 1$$

$$w_0 + \sum_i w_i X_i \rightarrow -\infty, P(Y = 1|X) \rightarrow 0$$

## Training Logistic Regression

Let's focus on binary classification

$$P(Y = 1|X) = \frac{\exp(w_0 + \sum_i w_i X_i)}{\exp(w_0 + \sum_i w_i X_i) + 1}$$

$$P(Y = 1|X) = \frac{1}{\exp(w_0 + \sum_i w_i X_i) + 1}$$

How to learn  $w_0, w_1 \dots w_d$ ?

Training data:  $\{(X^{(j)}, Y^{(j)})\}_{j=1}^n$ , with  $X^{(j)} = (X_1^{(j)}, X_2^{(j)}, \dots, X_d^{(j)})$

Maximum Likelihood Estimation:

$$\hat{\mathbf{w}}_{\text{MLE}} = \underset{\mathbf{w}}{\operatorname{argmax}} \prod_{j=1}^n P(X^{(j)}, Y^{(j)} | \mathbf{w})$$

## Optimizing concave/convex function

- $l(\mathbf{w})$  concave, we can maximize it via gradient ascent

Gradient:

$$\nabla_{\mathbf{w}} l(\mathbf{w}) = \left[ \frac{\partial l(\mathbf{w})}{\partial w_0}, \dots, \frac{\partial l(\mathbf{w})}{\partial w_d} \right]$$

Update rule for gradient ascent, with **learning rate  $\eta > 0$**

$$\Delta \mathbf{w} = \eta \nabla_{\mathbf{w}} l(\mathbf{w})$$

$$w_i^{(t+1)} = w_i^{(t)} + \eta \frac{\partial l(\mathbf{w})}{\partial w_i} \Big|_{w_t}$$

# Optimizing concave/convex function

- It's more common to use gradient descent to minimize a convex function

Update rule for gradient **descent**, with learning rate  $\eta > 0$

$$\Delta \mathbf{w} = -\eta \nabla_{\mathbf{w}} l(\mathbf{w})$$

$$w_i^{(t+1)} = w_i^{(t)} - \eta \frac{\partial l(\mathbf{w})}{\partial w_i} \Big|_{w_t}$$

(maximizing  $l(\mathbf{w})$  is the same as minimizing  $l'(\mathbf{w}) = -l(\mathbf{w})$ )

## Gradient Descent

Review, let's start with a simple function:

$$f(w) = 0.2(w - 2)^2 + 1$$

We know that this function is convex (2nd derivative exists and is positive).


```
In [5]: f = lambda w: 0.2*(w-2)**2+1  
        dfdw = lambda w: 0.4*w - 0.8
```

```
In [6]: w = np.linspace(-4,8,1000)
plt.plot(w, f(w), linewidth=3 )
plt.xlabel(r'$w$')
plt.ylabel(r'$f(w)$')
plt.title(r'Minimize $f(w)$, start with a random point $w_0$', fontsize = 20);
w_0 = 6
plt.plot(w_0, f(w_0), "o", markersize=10)

def draw_vector_2D(ax, x, y, lenx, leny, name, color='k'):
#     grad = np.array([-np.sin(x), np.cos(y)])
    ax.quiver(x,y,lenx, leny, color=color, angles='xy', scale_units='xy', scale=1)
    ax.text(x+lenx/2, y+leny/2+0.5, name, fontsize = 16, color=color)

draw_vector_2D(plt, w_0, 0, dfdw(w_0), 0, r'$\nabla f(w_0)$', 'k')
```

Minimize  $f(w)$ , start with a random point  $w_0$



```

In [7]: plt.plot(w, f(w), linewidth=3 )
plt.xlabel(r'$w$')
plt.ylabel(r'$f(w)$')

plt.title(r'Minimize $f(w)$, start with a random point $w_0$, step size $\eta=0.5$
$', fontsize = 20);
w_0 = 6
plt.plot(w_0, f(w_0), "o", markersize=10)


draw_vector_2D(plt, w_0, 0, dfdw(w_0), 0, r' ', 'k')

eta=0.5

draw_vector_2D(plt, w_0, 0, - dfdw(w_0)*eta, 0, r'$-\eta\nabla f(w_0)$', 'r')

```

Minimize  $f(w)$ , start with a random point  $w_0$ , step size  $\eta = 0.5$



```

In [8]: plt.plot(w, f(w), linewidth=3 )
plt.xlabel(r'$w$')
plt.ylabel(r'$f(w)$')


w_1 = w_0 - dfdw(w_0)*eta

plt.title(r'Set  $w_{t+1} = w_t - \eta \nabla f(w_t)$  and repeat', fontsize = 20
);

plt.plot(w_0, f(w_0), "o",markersize=10)
plt.plot(w_1, f(w_1), "o",markersize=10)

draw_vector_2D(plt, w_1, 0, - dfdw(w_1)*eta,0, r'$-\eta \nabla f(w_1)$', 'g')

```



```

In [9]: plt.plot(w, f(w), linewidth=3 )
plt.xlabel(r'$w$')
plt.ylabel(r'$f(w)$')

#  $w_1 = w_0 - dfdw(w_0) * \eta$ 
w_t = np.zeros(10)
w_t[0] = 7 #  $w_0$ 

eta = 4


for i in range(1,10):
    w_t[i] = w_t[i-1] - eta * dfdw(w_t[i-1] )

plt.title(r'Gradient descent with  $\eta=\{ }\$'.format(eta), fontsize = 20);

plt.plot(w_t, f(w_t), "o-", markersize=10)

# draw_vector_2D(plt, w_1, 0, - dfdw(w_1)*eta,0, r'$-\eta \nabla f(w_1)$', 'r')$ 
```

Out[9]: [



**Let's plot a function with two variables and look at the gradient**




```
In [10]: x = np.linspace(-1,2,100);y = np.linspace(-1,2,100); X,Y = np.meshgrid(x, y)

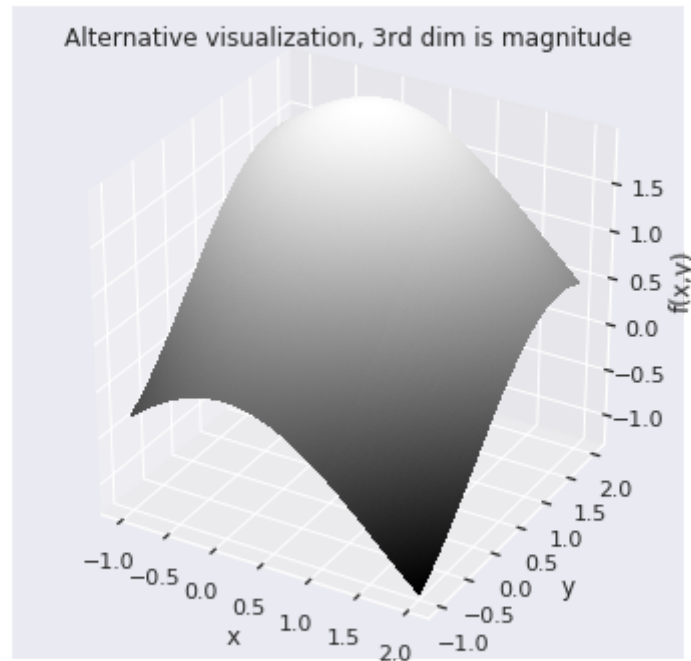
f_XY = np.cos(X)+np.sin(Y)

plt.figure(figsize=(6,5))
cs = plt.contourf(X, Y, f_XY,20,cmap='RdBu_r',vmin=-1,vmax=1,alpha=0.6);plt.colorbar()
ar()
contours = plt.contour(cs, colors='k')
plt.xlabel('x');plt.ylabel('y')
plt.title('color indicates magnitude of funciton \n gradient is perpendicular to l
evel set')

draw_vector_2D(plt, 1.45,0.5,-np.sin(1.45),np.cos(0.5),'','k')
```



```
In [11]: from mpl_toolkits.mplot3d import Axes3D
%matplotlib inline
%config InlineBackend.print_figure_kwargs = {'bbox_inches':None}
fig = plt.figure(figsize=(6,6))
ax = fig.gca(projection='3d')
Z = np.cos(X)+np.sin(Y)
# Plot the surface.
surf = ax.plot_surface(X, Y, Z, cmap='gray',
                      linewidth=0, antialiased=False, rcount=200, ccount=200)
ax.set_xlabel('x');ax.set_ylabel('y');ax.set_zlabel('f(x,y)');
ax.set_title('Alternative visualization, 3rd dim is magnitude');
```



# Logistic regression gradient ascent

Simple simulated example


```
In [12]: # Previous example
mu_1_1 = -5; sigma_1_1 = 2; mu_2_1 = 5; sigma_2_1 = 2
mu_1_0 = 5; sigma_1_0 = 2; mu_2_0 = -5; sigma_2_0 = 2
cov_positive = np.array([[sigma_1_1**2,3], [3,sigma_2_1**2]] )
cov_negative = np.array([[sigma_1_0**2,3], [3,sigma_2_0**2]] )
# Sample data from these distributions
X_positive = multivariate_normal.rvs(mean=[mu_1_1,mu_2_1], cov=cov_positive, size
= (20))
X_negative = multivariate_normal.rvs(mean=[mu_1_0,mu_2_0], cov=cov_negative, size
= (20))

X = np.vstack([X_positive, X_negative])
Y = np.vstack([np.ones((X_positive.shape[0],1)),np.zeros((X_negative.shape[0],1
))])
```

```
In [13]: plt.figure(figsize=(6,6))

plt.scatter(X_positive[:, 0], X_positive[:, 1],facecolors='r', edgecolors='w')
plt.scatter(X_negative[:, 0], X_negative[:, 1],facecolors='b', edgecolors='w')
plt.plot(x1, x1*0.8)

plt.axis([-10,10,-10,10],'equal')
plt.xlabel(r'$x_1$',fontsize=20)
plt.ylabel(r'$x_2$',fontsize=20)
plt.title('Data space',fontsize=20);
```



## Log likelihood plot

$$l(\mathbf{w}) = \sum_j \left[ y^j \left( w_0 + \sum_{i=1}^d w_i x_i^j \right) - \ln \left( 1 + \exp \left( w_0 + \sum_{i=1}^d w_i x_i^j \right) \right) \right]$$

We omit  $w_0$  in the example below for simplicity


```
In [14]: w1x = np.linspace(-60,40,100)
w2x = np.linspace(-60,40,100)
W1,W2 = np.meshgrid(w1x, w2x)

## ommiting w_0 just for illustration
def loglikelihood(w1,w2):
    w = np.array([[w1],[w2]]) # make w_vec
    loglikelihood = np.sum(Y*X.dot(w) - np.log(1+ np.exp(X.dot(w))))
    return loglikelihood

L_w = np.vectorize(loglikelihood)(*np.meshgrid(w1x, w2x))
```

```
In [15]: plt.figure(figsize=(8,6))

cs = plt.contourf(W1, W2, L_w,20,cmap='RdBu_r',vmin=-np.max(np.abs(L_w)),
                  vmax=np.max(np.abs(L_w)),alpha=0.6);
plt.colorbar()
contours = plt.contour(cs, colors='k')
plt.xlabel(r'$w_1$',fontsize=20)
plt.ylabel(r'$w_2$',fontsize=20)
plt.title('Parameter space',fontsize=20);
```



## Gradient computation

$$l(\mathbf{w}) = \sum_j \left[ y^j \left( w_0 + \sum_{i=1}^d w_i x_i^j \right) - \ln \left( 1 + \exp \left( w_0 + \sum_{i=1}^d w_i x_i^j \right) \right) \right]$$
$$\frac{\partial l(\mathbf{w})}{\partial w_0} = \sum_j \left[ y^j - \frac{\exp(w_0 + \sum_{i=1}^d w_i x_i^j)}{(1 + \exp(w_0 + \sum_{i=1}^d w_i x_i^j))} \right]$$
$$= \sum_j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})]$$

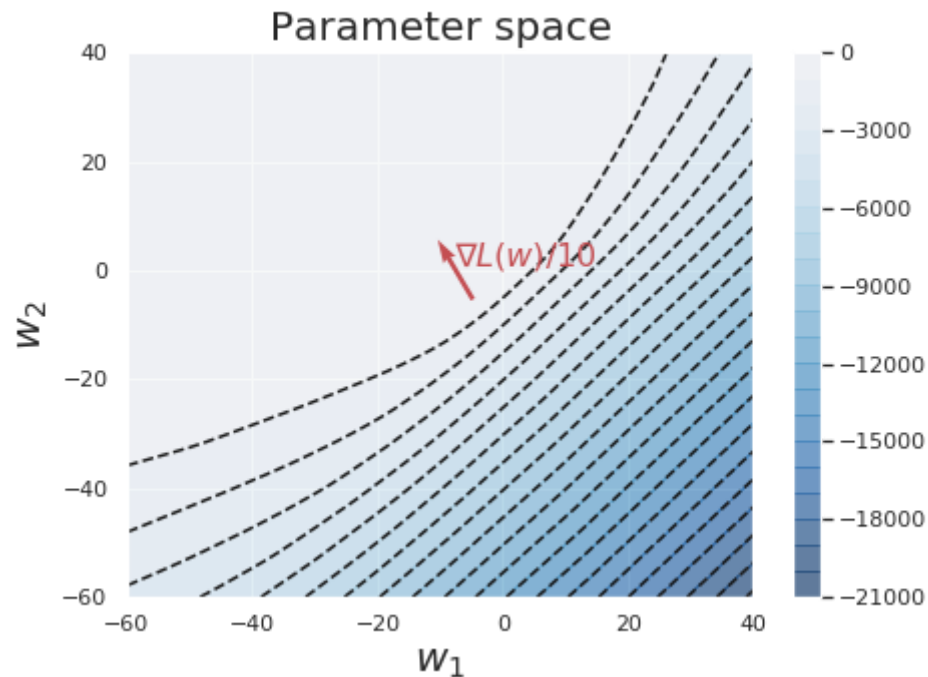
for  $i = 1 \dots d$ :

$$\frac{\partial l(\mathbf{w})}{\partial w_0} = \sum_j x_i^j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})]$$

```
In [16]: def gradient_likelihood(w1,w2,X,Y):  
    w = np.array([[w1],[w2]])  
    P_Y_1 = np.exp(X.dot(w))/(1+ np.exp(X.dot(w)))  
    gw1 = X[:,0:1].T.dot(Y-P_Y_1)  
    gw2 = X[:,1:2].T.dot(Y-P_Y_1)  
    return gw1, gw2
```

```
In [17]: plt.figure(figsize=(7,5))
cs = plt.contourf(W1, W2, L_w,20,cmap='RdBu_r',vmin=-np.max(np.abs(L_w)),
                  vmax=np.max(np.abs(L_w)),alpha=0.6); plt.colorbar()
contours = plt.contour(cs, colors='k')
plt.xlabel(r'$w_1$',fontsize=20);plt.ylabel(r'$w_2$',fontsize=20)
plt.title('Parameter space',fontsize=20);

w1 = -5; w2 = -5
gw1, gw2 = gradient_likelihood(w1,w2,X, Y)
draw_vector_2D(plt, w1,w2,gw1/10,gw2/10, r'$\nabla L(w)/10$', 'r');
```





# Gradient ascent for logistic regression

Iterate until convergence (until change  $< \epsilon$ )

$$w_0^{(t+1)} = w_0^{(t)} + \eta \sum_j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})]$$

for  $i = 1 \dots d$ :

$$w_i^{(t+1)} = w_i^{(t)} + \eta \sum_j x_i^j [y^j - \hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})]$$

$\hat{P}(Y^j = 1 | \mathbf{x}^j, \mathbf{w}^{(t)})$  is our current prediction of the label.

- compare this to actual label
- multiple difference by feature value

# Gradient Ascent (Descent) is simplest of approaches

Compare to:

- Stochastic Gradient Descent
- Batch Gradient Descent
- Newton

...

- Convergence?

Effect of step-size  $\eta$  Large  $\eta \Rightarrow$  Fast convergence but larger residual error Also possible oscillations Small  $\eta \Rightarrow$  Slow convergence but small residual error

## Need to regularize the weights

- $w \rightarrow \infty$  if the data is linearly separable
- For MAP, need to define prior on  $W$ 
  - given  $W = (w_1, \dots, w_d)$
  - let's assume prior  $P(w_i) = \mathcal{N}(0, \sigma)$
- A kind of Occam's razor (simplest is best) prior
- Helps avoid very large weights and overfitting

## Adding a prior on $W$

MAP estimation picks the parameter  $W$  that has maximum posterior probability  $P(W|Y, X)$  given the conditional likelihood  $P(Y|W, X)$  and the prior  $P(W)$ .

Using Bayes rule again:

$$\begin{aligned} W^{MAP} &= \operatorname{argmax}_W P(W|Y, X) = \operatorname{argmax}_W \frac{P(Y|W, X)P(W, X)}{P(Y, X)} \\ &= \operatorname{argmax}_W P(Y|W, X)P(W, X) \\ &= \operatorname{argmax}_W P(Y|W, X)P(W)P(X) \quad \text{assume } P(W, X) = P(W)P(X) \\ &= \operatorname{argmax}_W P(Y|W, X)P(W) \\ &= \operatorname{argmax}_W \ln P(Y|W, X) + \ln P(W) \end{aligned}$$

Zero Mean Gaussian prior on  $W$ :  $W \sim \frac{1}{2\pi\sigma^2} \exp \left( -\frac{1}{2\sigma^2} \sum_i w_i^2 \right)$


$$W^{MAP} = \operatorname{argmax}_W \ln P(Y|W, X) - \left( \frac{1}{2\sigma^2} \sum_i w_i^2 \right)$$

```
In [18]: lambda = 10 # this is 1/(2*sigma**2)

def logposterior(w1,w2):
    w = np.array([[w1],[w2]]) # make w_vec
    loglikelihood = np.sum(Y*X.dot(w) - np.log(1+ np.exp(X.dot(w))))
    loglikelihood += - (w1**2 + w2**2)*lambda
    return loglikelihood

L_w = np.vectorize(logposterior)(*np.meshgrid(w1x, w2x))
```


```
In [19]: plt.figure(figsize=(8,6))
cs = plt.contourf(W1, W2, L_w,20,cmap='RdBu_r',vmin=-np.max(np.abs(L_w)),
                  vmax=np.max(np.abs(L_w)),alpha=0.6);
plt.colorbar()
contours = plt.contour(cs, colors='k')
plt.xlabel(r'$w_1$',fontsize=20)
plt.ylabel(r'$w_2$',fontsize=20)
plt.title('Parameter space',fontsize=20);
```



```
In [20]: def gradient_posterior(w1,w2,X,Y):  
    w = np.array([[w1],[w2]])  
    P_Y_1 = np.exp(X.dot(w))/(1+ np.exp(X.dot(w)))  
    gw1 = X[:,0:1].T.dot(Y-P_Y_1)- 2*lmbda*w1#  
  
    gw2 =X[:,1:2].T.dot(Y-P_Y_1) - 2*lmbda*w2 #  
  
    return gw1, gw2
```

```
In [22]: plt.figure(figsize=(8,6))
cs = plt.contourf(W1, W2, L_w,20,cmap='RdBu_r',vmin=-np.max(np.abs(L_w)),
                  vmax=np.max(np.abs(L_w)),alpha=0.6); plt.colorbar()
contours = plt.contour(cs, colors='k')
plt.xlabel(r'$w_1$',fontsize=20);plt.ylabel(r'$w_2$',fontsize=20)
plt.title('Parameter space',fontsize=20);

w1 = 7; w2 = -16
gw1, gw2 = gradient_likelihood(w1,w2,X, Y)
draw_vector_2D(plt, w1,w2,gw1/10,gw2/10, r' ', 'r');
gw1, gw2 = gradient_posterior(w1,w2,X, Y)
draw_vector_2D(plt, w1,w2,gw1/10,gw2/10, r'$\nabla L(w)/10$', 'b');
```





## What you should know

LR is a linear classifier: decision rule is a hyperplane

- LR optimized by conditional likelihood
  - no closed-form solution
  - concave  $\Rightarrow$  global optimum with gradient ascent
  - Maximum conditional a posteriori corresponds to regularization

In [ ]: